

Word meaning as classification

CMSC 723 / LING 723 / INST 725

Hal Daumé III [he/him]

17 Sep 2019

(Many slides c/o Dan Jurafsky & James Martin)

Announcements, logistics

- HW2 (ddl Sept 26 before class):
 - Written portion was posted last week
 - You should have received a login for Jupyter Hub server via email
if not please contact us ASAP (either after class to talk to Amr)
 - Programming portion posted in next 24 hours (hopefully today)!

- Readings for next few weeks are now posted

R 19 Sep	Data collection and annotation	DataInNLP , and AnnCaseStudy	
T 24 Sep	Measurement and validity	Measurement , and MeasurementCaseStudy , Sec "Reliability, Validity, ..."	
R 26 Sep	Crowdsourcing annotations	CrowdsourcingNLP , and AnnMyths	HW2
T 01 Oct	Multilinguality and linguistic variety	TheBenderRule , and Elicitation , Sec 3, and optional: ActiveElicitation	

Tf-idf and PPMI are
sparse representations

- **tf-idf and PPMI vectors are**
 - **long** (length $|V| = 20,000$ to $50,000$)
 - **sparse** (most elements are zero)

Alternative: dense vectors

- vectors which are
 - **short** (length 50-1000)
 - **dense** (most elements are non-zero)

Sparse versus dense vectors

- Why dense vectors?
 - Short vectors may be easier to use as **features** in machine learning (less weights to tune)
 - Dense vectors may **generalize** better than storing explicit counts
 - They may do better at capturing synonymy:
 - *car* and *automobile* are synonyms; but are distinct dimensions
 - a word with *car* as a neighbor and a word with *automobile* as a neighbor should be similar, but aren't
 - **In practice, they work better**

Dense embeddings you can download!

- **word2vec**

- <https://code.google.com/archive/p/word2vec/>

- **Fasttext**

- <http://www.fasttext.cc/>

- **Glove**

- <http://nlp.stanford.edu/projects/glove/>

Word2vec

- Popular embedding method
- Very fast to train
- Code available on the web
- Idea: **predict** rather than **count**

Word2vec

- Instead of **counting** how often each word w occurs near "*apricot*"
- Train a classifier on a binary **prediction** task:
 - Is w likely to show up near "*apricot*"?
- We don't actually care about this task
 - But we'll take the learned classifier weights as the word embeddings

Use running text as implicitly supervised training data!

- A word s near *apricot*
 - Acts as gold ‘correct answer’ to the question
 - “Is word w likely to show up near *apricot*?”
- No need for hand-labeled supervision

Word2Vec: Skip-Gram Task

- Word2vec provides a variety of options
 - Let's do "skip-gram with negative sampling" (SGNS)

Skip-gram algorithm

1. Treat the target word and a neighboring context word as positive examples.
2. Randomly sample other words in the lexicon to get negative samples
3. Use logistic regression to train a classifier to distinguish those two cases
4. Use the weights as the embeddings

Skip-Gram Training Data

- Training sentence:
- ... lemon, a tablespoon of apricot jam a pinch ...
- c1 c2 target c3 c4

Assume context words are those in +/- 2
word window

Skip-Gram Goal

- Given a tuple (t,c) = target, context
 - (*apricot*, *jam*)
 - (*apricot*, *aardvark*)
- Return probability that c is a real context word:
 - $P(+ | t,c)$
 - $P(- | t,c) = 1 - P(+ | t,c)$

How to compute $p(+ | t, c)$?

- Intuition:
 - Words are likely to appear near similar words
 - Model similarity with dot-product!
- Classification model
 - $P(+ | t, c) = 1/(1+\exp(-e_t \cdot e_c))$

Skip-Gram Training Data

- Training sentence:
- ... lemon, a tablespoon of apricot jam a pinch ...
- c1 c2 t c3 c4
- Training data: input/output pairs centering on *apricot*
- Assume a +/- 2 word window

Skip-Gram Training

- Training sentence:

- ... lemon, a tablespoon of apricot jam a pinch

...

- | | | | | | | | |
|--|----|--|----|---|--|----|----|
| | c1 | | c2 | t | | c3 | c4 |
|--|----|--|----|---|--|----|----|

positive examples +

t	c
<hr/>	
apricot	tablespoon
apricot	of
apricot	preserves
apricot	or

- For each positive example, we'll create k negative examples.

- Using *noise* words

- Any random word that isn't t

16

Skip-Gram Training

- Training sentence:

- ... lemon, a tablespoon of apricot jam a pinch

...

- c1 c2 t c3 c4

positive examples +

t	c
---	---

apricot	tablespoon
---------	------------

apricot	of
---------	----

apricot	preserves
---------	-----------

apricot	or
---------	----

negative examples - ^{k=2}

t	c	t	c
---	---	---	---

apricot	aardvark	apricot	twelve
---------	----------	---------	--------

apricot	puddle	apricot	hello
---------	--------	---------	-------

apricot	where	apricot	dear
---------	-------	---------	------

apricot	coaxial	apricot	forever
---------	---------	---------	---------

¹⁷

Choosing noise words

- Could pick w according to their unigram frequency $P(w)$
- More common to chosen then according to $p_\alpha(w)$

$$P_\alpha(w) = \frac{\text{count}(w)^\alpha}{\sum_w \text{count}(w)^\alpha}$$

- $\alpha = \frac{3}{4}$ works well because it gives rare noise words slightly higher probability
- To show this, imagine two events $p(a) = .99$ and $p(b) = .01$:

$$P_\alpha(a) = \frac{.99^{.75}}{.99^{.75} + .01^{.75}} = .97$$

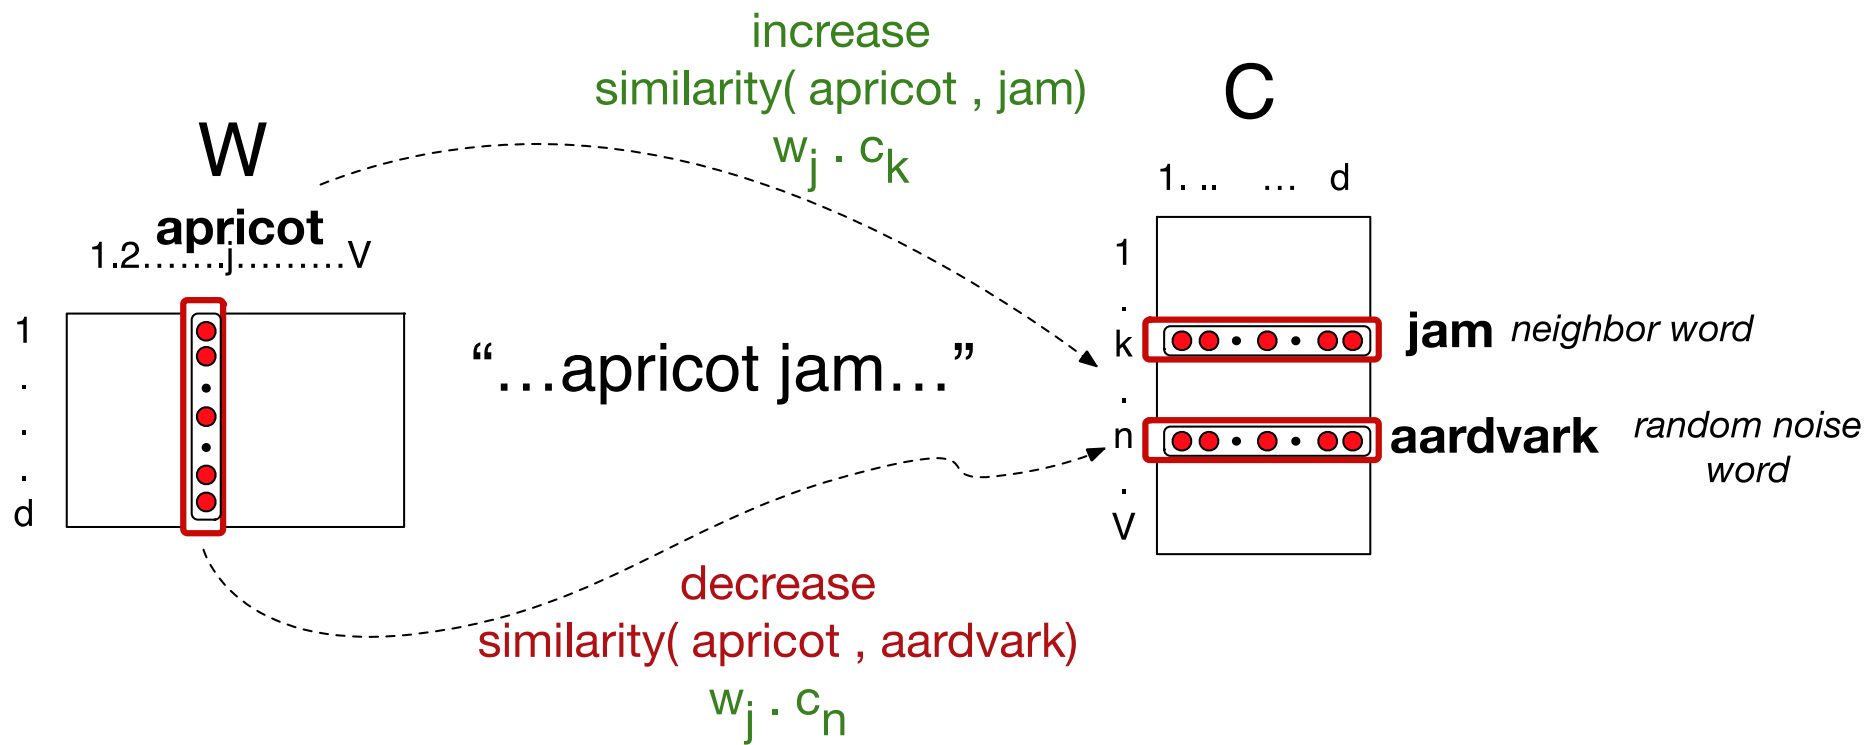
$$P_\alpha(b) = \frac{.01^{.75}}{.99^{.75} + .01^{.75}} = .03$$

Setup

- Let's represent words as vectors of some length (say 300), randomly initialized.
- So we start with $300 * V$ random parameters
- Over the entire training set, we'd like to adjust those word vectors such that we
 - Maximize the similarity of the **target word**, **context word** pairs (t,c) drawn from the positive data
 - Minimize the similarity of the (t,c) pairs drawn from the negative data.

Learning the classifier

- Iterative process.
- We'll start with 0 or random weights
- Then adjust the word weights to
 - make the positive pairs more likely
 - and the negative pairs less likely
- over the entire training set:



Train using gradient descent

- Actually learns two separate embedding matrices W and C
- Can use W and throw away C , or merge them somehow

Summary: How to learn word2vec (skip-gram) embeddings

- Start with V random 300-dimensional vectors as initial embeddings
- Use logistic regression, the second most basic classifier used in machine learning after naïve bayes
 - Take a corpus and take pairs of words that co-occur as positive examples
 - Take pairs of words that don't co-occur as negative examples
 - Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
 - Throw away the classifier code and keep the embeddings.

Evaluating embeddings

- Compare to human scores on word similarity-type tasks:
 - WordSim-353 (Finkelstein et al., 2002)
 - SimLex-999 (Hill et al., 2015)
 - Stanford Contextual Word Similarity (SCWS) dataset (Huang et al., 2012)
 - TOEFL dataset: *Levied is closest in meaning to: imposed, believed, requested, correlated*

Properties of embeddings


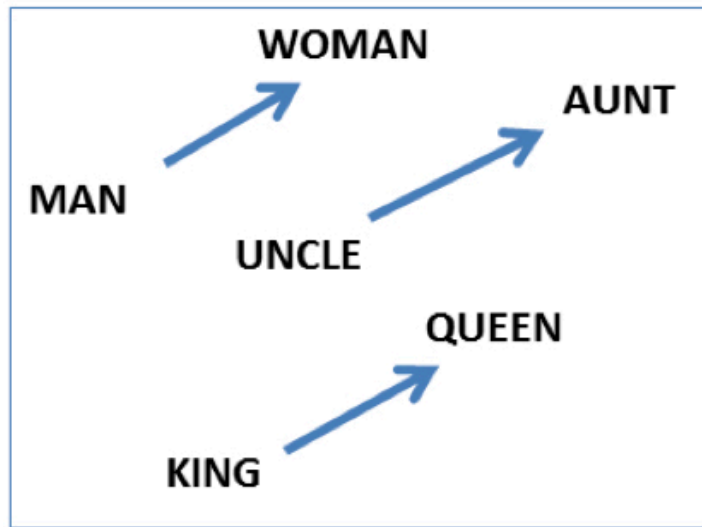
Similarity depends on window size C


- $C = \pm 2$ The nearest words to *Hogwarts*:
 - *Sunnydale*
 - *Evernight*
- $C = \pm 5$ The nearest words to *Hogwarts*:
 - *Dumbledore*
 - *Malfoy*
 - *halfblood*


Analogy: Embeddings capture relational meaning!

$\text{vector}('king') - \text{vector}('man') + \text{vector}('woman') \approx \text{vector}('queen')$

$\text{vector}('Paris') - \text{vector}('France') + \text{vector}('Italy') \approx \text{vector}('Rome')$

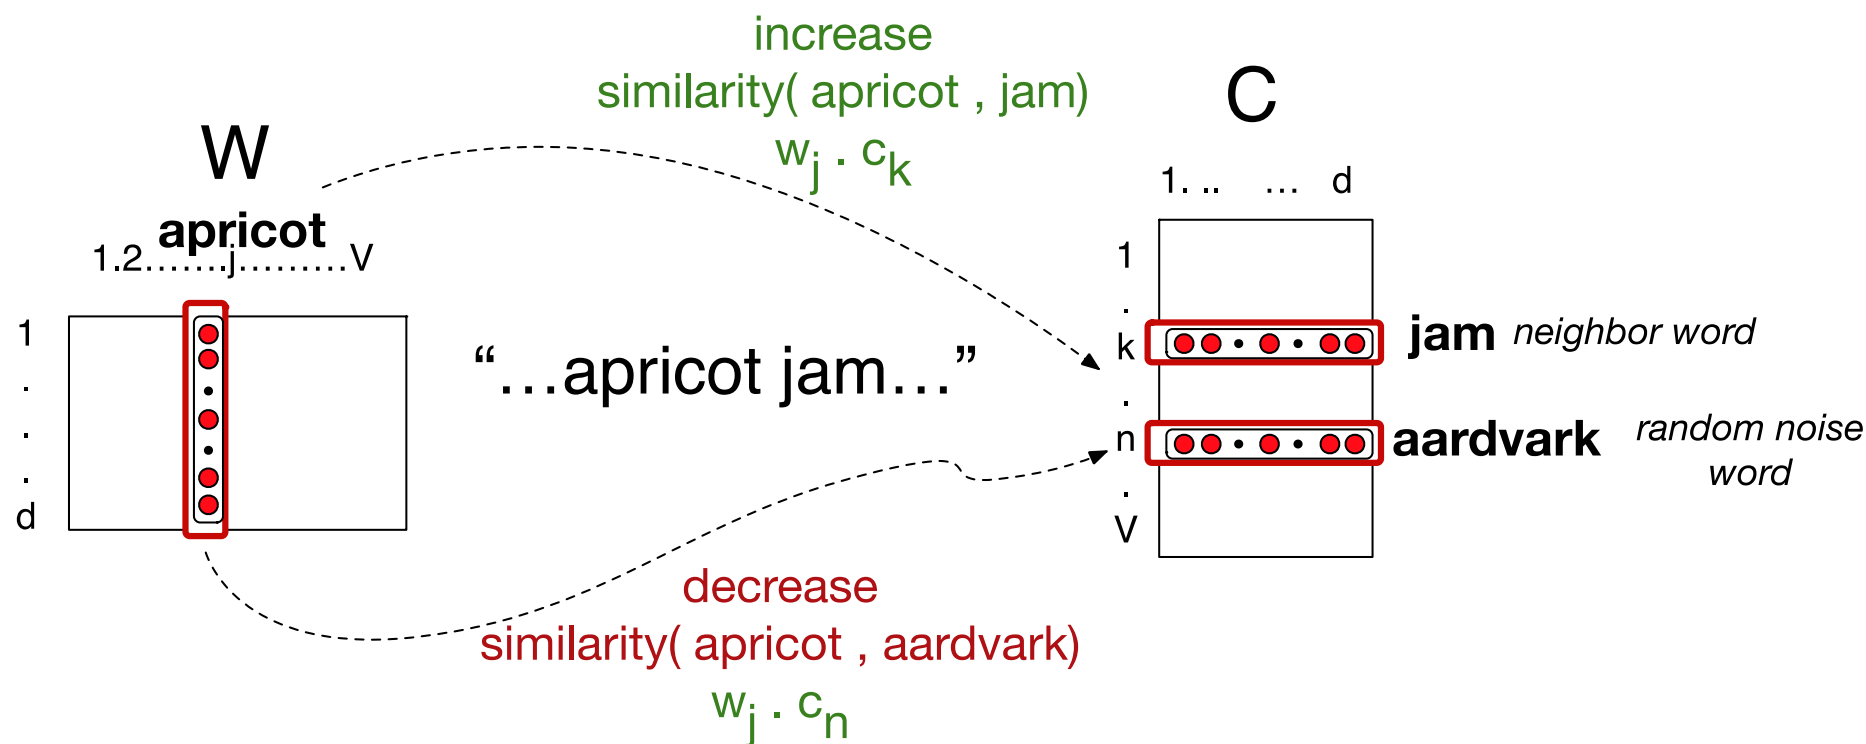






Are embeddings so magical?


Word2vec as matrix factorization



$M = WC$, but... what is M ? (Levy & Goldberg, NeurIPS 2014)

$$M = \log \left(\frac{\#(w, c) \cdot |D|}{\#(w) \cdot \#(c)} \right) - \log k = PMI(w_i, c_j) - \log k$$

How does
this help us
understand
these
“semantic”
embedding
dimensions?



Does this matter empirically?

WS353 (WORDSIM) [13]			MEN (WORDSIM) [4]			MIXED ANALOGIES [20]		SYNT. ANALOGIES [22]			
Representation		Corr.	Representation		Corr.	Representation		Acc.	Representation		Acc.
SVD	(k=5)	0.691	SVD	(k=1)	0.735	SPPMI	(k=1)	0.655	SGNS	(k=15)	0.627
SPPMI	(k=15)	0.687	SVD	(k=5)	0.734	SPPMI	(k=5)	0.644	SGNS	(k=5)	0.619
SPPMI	(k=5)	0.670	SPPMI	(k=5)	0.721	SGNS	(k=15)	0.619	SGNS	(k=1)	0.59
SGNS	(k=15)	0.666	SPPMI	(k=15)	0.719	SGNS	(k=5)	0.616	SPPMI	(k=5)	0.466
SVD	(k=15)	0.661	SGNS	(k=15)	0.716	SPPMI	(k=15)	0.571	SVD	(k=1)	0.448
SVD	(k=1)	0.652	SGNS	(k=5)	0.708	SVD	(k=1)	0.567	SPPMI	(k=1)	0.445
SGNS	(k=5)	0.644	SVD	(k=15)	0.694	SGNS	(k=1)	0.540	SPPMI	(k=15)	0.353
SGNS	(k=1)	0.633	SGNS	(k=1)	0.690	SVD	(k=5)	0.472	SVD	(k=5)	0.337
SPPMI	(k=1)	0.605	SPPMI	(k=1)	0.688	SVD	(k=15)	0.341	SVD	(k=15)	0.208

Where did things go from here...

- Training objectives (fasttext, glove, etc.)
- Multilingual embeddings (also fasttext, others)
- Lots of “bias in embeddings” stuff
- Trying to tease out what embeddings are actually learning
- Using embeddings as an input representation is a very good default for dense models. But bag of n-grams still often wins

Side-note on biases

- There are >64 papers on biases and “debiasing”
- I have many thoughts on this, we’ll talk about it later
- For now:
 - Embeddings reflect co-occurrence in data
 - Co-occurrences reflect social structure as reflected in text
 - Text != the real world

Where did things go from here...

- Training objectives (fasttext, glove, etc.)
- Multilingual embeddings (also fasttext, others)
- Lots of “bias in embeddings” stuff
- Trying to tease out what embeddings are actually learning
- Using embeddings as an input representation is a very good default for dense models. But bag of n-grams still often wins

Today

- Embeddings
- Word2vec
- Word2vec isn't magic